

Large-scale cluster management at Google with Borg

By: Abhishek Verma, Luis Pedrosa, Madhukar Korupolu, David Oppenheimer, Eric Tune, John Wilkes

Papers We Love - Brasília

December 15th, 2016

What is Borg?

- A cluster management system developed by Google

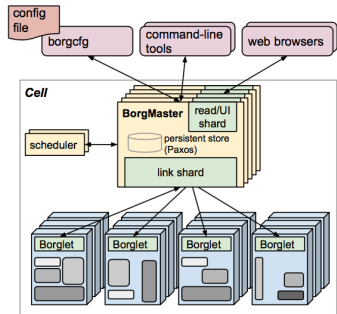


Figure 1: Borg's high-level architecture

What is Borg?

- A cluster management system developed by Google
- Its motivations are:

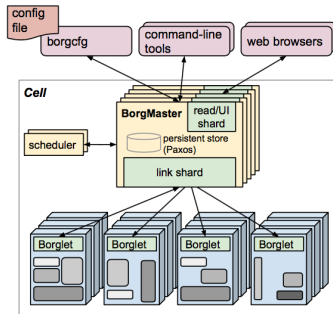


Figure 1: Borg's high-level architecture

What is Borg?

- A cluster management system developed by Google
- Its motivations are:
 - 1 hide the details of resource management from the users. In other words, to help the software engineers to focus on writing code, instead of dealing with application deployment




Figure 1: Borg's high-level architecture

What is Borg?

- A cluster management system developed by Google
- Its motivations are:
 - 1 hide the details of resource management from the users. In other words, to help the software engineers to focus on writing code, instead of dealing with application deployment
 - 2 provide high reliability and availability environment for the applications

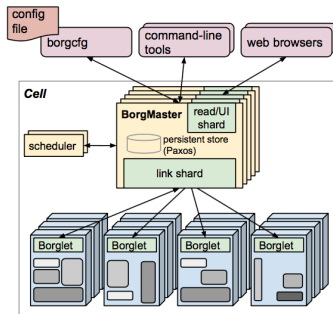


Figure 1: Borg's high-level architecture

What is Borg?

- A cluster management system developed by Google
- Its motivations are:
 - 1 hide the details of resource management from the users. In other words, to help the software engineers to focus on writing code, instead of dealing with application deployment
 - 2 provide high reliability and availability environment for the applications
 - 3 improve resource sharing




Figure 1: Borg's high-level architecture

The users are mainly software engineers and site reliability engineers (SRE)

- Borg's workload comprises:

The users are mainly software engineers and site reliability engineers (SRE)

- Borg's workload comprises:
 - 1 long-running services: should “never” go down

- Borg's workload comprises:
 - 1 long-running services: should “never” go down
 - 2 batch jobs: take few seconds to few days to complete

- Borg's workload comprises:
 - ① long-running services: should “never” go down
 - ② batch jobs: take few seconds to few days to complete
- Borg provides a DSL to allow the users to describe the requirement for their jobs

```
job hello_world = {  
  runtime = { cell = 'ic' }           // What cluster should we run in?  
  binary = '../hello_world_webserver' // What program are we to run?  
  args = { port = '%port%' }         // Command line parameters  
  requirements = {                   // Resource requirements  
    ram = 100M  
    disk = 100M  
    cpu = 0.1  
  }  
  replicas = 10000 // Number of tasks  
}
```

- Borg's workload comprises:
 - ① long-running services: should “never” go down
 - ② batch jobs: take few seconds to few days to complete
- Borg provides a DSL to allow the users to describe the requirement for their jobs

```
job hello_world = {  
  runtime = { cell = 'ic' }           // What cluster should we run in?  
  binary = '../hello_world_webserver' // What program are we to run?  
  args = { port = '%port%' }         // Command line parameters  
  requirements = {                   // Resource requirements  
    ram = 100M  
    disk = 100M  
    cpu = 0.1  
  }  
  replicas = 10000 // Number of tasks  
}
```

- Likewise, the jobs are classified as production (i.e., higher-priority) and non-production

The background image shows a vast data center with a complex steel truss ceiling and rows of server racks. A semi-transparent blue box with white text is overlaid in the center. At the bottom of the image, there is a navigation bar with various icons.

Cells

- ❶ cluster of machines located in a single data center building
- ❷ a median cluster size is 10K
- ❸ the machines are heterogeneous (CPU, RAM, disk, network)
- ❹ Cells are connected by high-speed networks

1 Allocs

- Reserved set of resources

2 Priority, quota, and admission control

- job has a priority
- Quota is used to decide which jobs to admit for scheduling

3 Naming and Monitoring

- 50.jfoo.ubar.cc.borg.google.com
- Monitoring health of the task and thousands of performance metrics

Borg follows a master/slave architecture

1 Borgmaster

- Borgmaster process and scheduler
- Replicated

2 Borglet

- Manage and monitor tasks and resource
- Borgmaster polls Borglet every few seconds

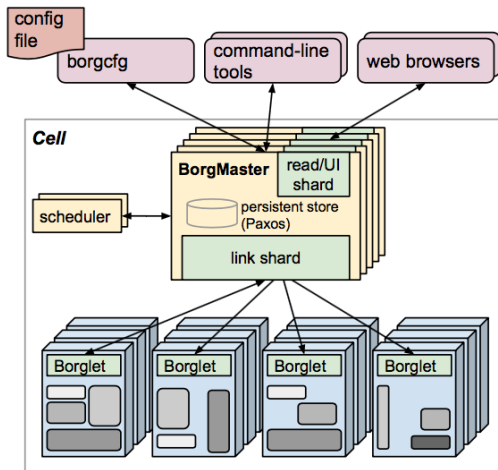


Figure 2: Borg's high-level architecture

A scoring system combines user-specified preferences & build-in criteria

1 Scheduling

- **feasibility checking:** find machines on which the task could run
- **scoring:** pick a machine considering different criteria

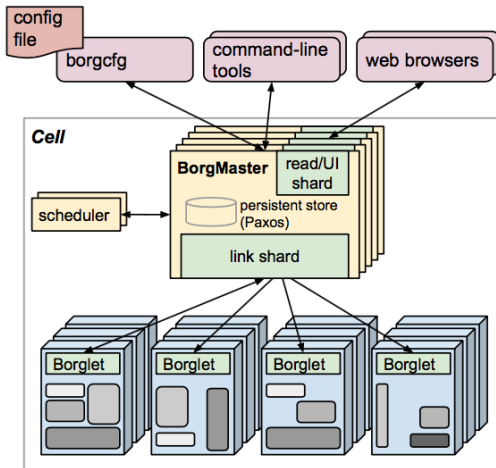


Figure 3: Borg's high-level architecture ↻ 🔍 🔄

Borg employs a modular approach to run its functions across different processes

- 1 Separate scheduler
- 2 Separate threads to poll the Borglets
- 3 Partition functions across the five replicas
- 4 Score caching: Cache score, ignore small changes
- 5 Equivalence classes: scoring for one task per equivalence class
- 6 Relaxed randomization: don't scores all machines, choose random number of machines

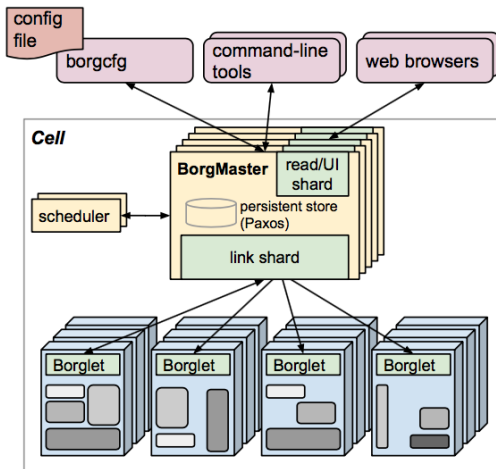




Figure 4: Borg's high-level architecture


Availability




- 1 Implements checkpoint
- 2 Reschedules evicted tasks
- 3 Spreads tasks across failure domain



- 1 **Cell sharing:** segregating prod and non-prod works into different cells would need more machines




(a) The left column for each cell shows the original size and the combined workload; the right one shows the segregated case.




(b) CDF of additional machines that would be needed if we segregated the workload of 15 representative cells.

- 2 **Cell size:** subdividing cells into smaller ones would also require more resources




(a) Additional machines that would be needed as a function of the number of smaller cells for five different original cells.



(b) A CDF of additional machines that would be needed divide each of 15 different cells into 2, 5 or 10 cells.

- 3 **Resource reclamation:** estimate how many resources a task will use and reclaim the rest for work

Utilization III



① Bad decisions

- jobs are restrictive as the only grouping mechanism for tasks
- one IP address per machine complicates things
- optimizing for power users at the expense of casual ones

② Good decisions

- use of *allocs*
- cluster management is more than task management
- introspection is primordial
- the master is the kernel of a distributed system

Kubernetes evolved from Borg

- 1 An open-source system for automating deployment, operations, and scaling of containerized applications
- 2 Pods: groups of containers
- 3 Labels
- 4 Replica controller
- 5 Services



That's all Folks!

